

# Chipyard Configurator

---

A RISC-V PARAMETERIZATION TOOL

Konstantinos Pikoulas  
Georgios Klempetsanis



## *What is RISC-V ?*

RISC-V is an ISA whose development began in 2010 at the University of California, Berkeley.

## *Why do we care?*

Because it is **open standard**!! RISC-V is offered under royalty-free open-source licenses.



## *What is Chipyard ?*

Chipyard is a framework composed of a collection of tools and libraries designed to provide an integration between open-source and commercial tools for the development of systems-on-chip.

## *Why are we using it?*


We are using Generators. A Generator can be thought of as a generalized RTL design, written using a mix of meta-programming and standard RTL. This type of meta-programming is enabled by the Chisel hardware description language.



# *But chisel code is so hard to maintain!*

That's why we made a python wrapper that generates chisel code for you!

## *How it works...*




*How many cores?  
L1/L2 cache configuration?  
Virtual Addressing?*

*Is it a power of 2 (e.g. cache ways)?  
Is the configuration valid?*

*Generate the chisel code based on  
the validated configuration.*

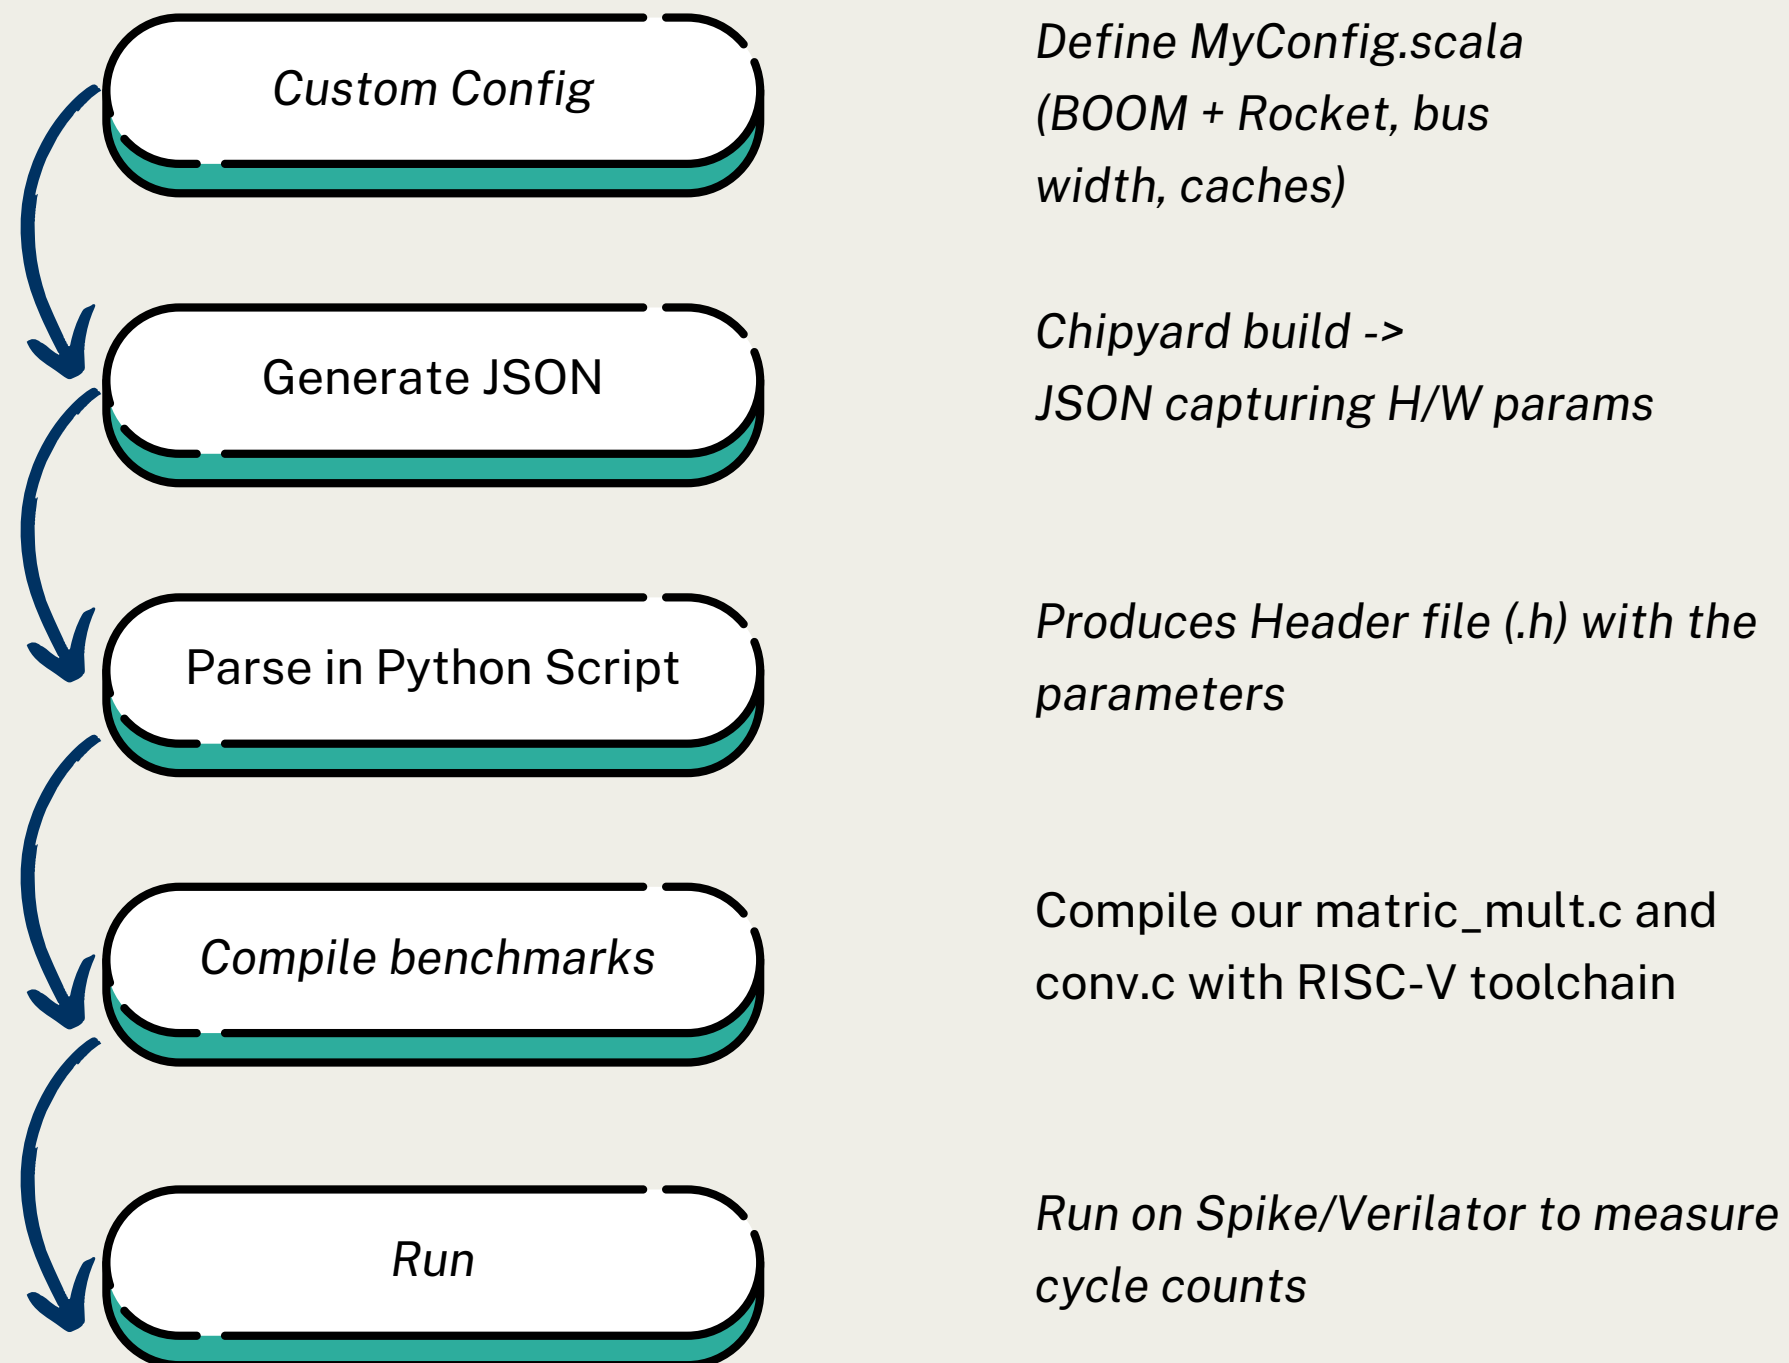
*Build the RTL.*



# Project Flow: From Config to Benchmark



## How it works...



## *Extracting Config Data from JSON*

- *We search the JSON for “cpus”, “cache sizes”, “bus widths”, etc*
- *Our script can adapt any fields we define in the Chipyard config*
- *Output macros, e.g:*
  1. *TOTAL\_CORES = sum of all CPU nodes*
  2. *L1\_DCACHE\_SIZE from “cache-sets \* cache-block-size \* ways”*
  3. *Additional data: system bus width, pipeline features etc*



JSON	Raw Data	Headers
Save	Copy	Collapse AllExpand AllFilter JSON
▶ #size-cells: [ ... ]		
▼ cpu@0:		
▶ clock-frequency: [ ... ]		
▼ compatible:		
0: "sifive,rocket0"		
1: "riscv"		
▼ d-cache-block-size:		
0: 64		
▼ d-cache-sets:		
0: 64		
▼ d-cache-size:		
0: 32768		
▼ d-tlb-sets:		
0: 1		
▼ d-tlb-size:		
0: 32		
▼ device_type:		
0: "cpu"		
▶ hardware-exec-breakpoint-count: [ ... ]		
▶ i-cache-block-size: [ ... ]		
▶ i-cache-sets: [ ... ]		
▶ i-cache-size: [ ... ]		
▶ i-tlb-sets: [ ... ]		
▶ i-tlb-size: [ ... ]		
▶ interrupt-controller: { ... }		
▶ mmu-type: [ ... ]		
▶ next-level-cache: [ ... ]		
▶ reg: [ ... ]		
▶ riscv,isa: [ ... ]		
▶ riscv,pmpgranularity: [ ... ]		
▶ riscv,pmpregions: [ ... ]		
▶ status: [ ... ]		
▶ timebase-frequency: [ ... ]		
tlb-split: [ ]		
▼ cpu@1:		
▶ clock-frequency: [ ... ]		
▼ compatible:		
0: "ucb-bar,boom0"		
1: "riscv"		
▼ d-cache-block-size:		
0: 64		
▼ d-cache-sets:		
0: 64		

# Config-Driven Loop Transformations

- *We define macros in code:*
  1. *#define TILE\_SIZE (...)*
  2. *#define UNROLL\_FACTOR (...)*
- *These macros compute from TOTAL\_CORES and L1\_DCACHE\_SIZE (or other parameters)*
- *Loop tiling, unrolling, merging, etc. become “plug-and-play”*





## Applying Tiling + Unrolling



- *Tiling:*

1. *Outer loops break the matrix dimension  $M$  into sub-blocks*

- *Unrolling:*

2. *Inner loop factor chosen from `TOTAL_CORES` or other performance heuristics*

- *If SoC has bigger caches  $\rightarrow$  bigger tile blocks*
- *If SoC has more cores  $\rightarrow$  higher unroll factor*

## 2D Tiling + Partial Kernel Unroll

- *For an  $N \times N$  input, if the SoC has bigger L1, we do bigger tile chunking over  $(i, j)$*
- *If `TOTAL_CORES` is large, we might unroll the kernel loops more aggressively*
- *Code picks these thresholds from macros; easy to extend for new SoCs*







# Collecting Cycle Counts on Arbitrary SoCs

- *Each new Chipyard config → Rebuild the SoC + produce new JSON*
- *Our script regenerates config\_params.h*
- *Recompile benchmarks → run on spike/verilator/FPGA*
- *Record cycle counts or other metrics (like instructions, energy, etc.)*



# Plotting the Results

- Bar plots / line charts vs. matrix size, kernel size, tile/unroll combos
- Identify best transformations for each SoC design





## *Key Takeaways (Generalized)*

- *Loop transformations remain crucial for large problem sizes and varied SoC designs*
- *Cache + core combos drive tiling/unrolling decisions*
- *Gains are currently single-threaded — true multicore parallelization is a next step*
- *Auto-tuning loop parameters or HPC frameworks is possible*
- *Any SoC config integrated with minimal changes to the pipeline*

## *Wrapping Up: A Universal Approach*

- *We built an open pipeline where any Chipyard config → auto param → optimized code*
- *Extensible to more loops or HPC libraries*
- *Next: multi-thread concurrency, advanced auto-tuning, or hardware-specific instructions*